

IMPLEMENTING PARALLELISM IN BFS

Kaushik Naraynan Saaketh Balachendil

March 26, 2024

Language used

C++

Libraries used

OpenMP

Algorithm

- Instead of popping out one vertex at a time, pop out all the nodes in the same level. (These nodes are known as frontier nodes)
- Level synchronous traversal. Each the processor will take a set of frontier vertices and calculate their next frontier vertices in parallel.
- For the above step we will need to partition the adjacency matrix and the vertices and allocate them to the processors.
- The adj matrix is divided into P blocks of size $N/\sqrt{p} \times N/\sqrt{p}$.
- Vertex are partitioned into N/P size groups.
- $N - i$ number of processors
- $p - i$ number of threads (for our case it is 4).
- Do a transpose of the frontier vector between the processors.
- After this all the columns processor s will have matching frontier with their local adjacency matrix.
- We then do a column wise all gather for the frontier vertices.
- This will broadcast the required frontier vertices for each column.
- Calculation of next frontier vertices is based on the current frontier vector that the processor has.

- Using the local adj matrix the next frontier vector is calculated.
- Note that each processor row now has the full information of the next frontier vertices.
- Now we do a all to all gather row wise so that all the next frontier vectors are merged. (union)
- All the processors now know if they have any frontier element (Next frontier now becomes current local frontier) that they own.
- We mark the node as visited and store its parent node.
- We do a row wise all gather and then column wise all gather to broadcast the local frontiers globally.
- We continue the process till there is no vertices left in the global frontier vertices.

Code

```

1  /*
2  * BFS Parallelization for undirected graphs
3  * Uses 4 threads to parallelize the traversal
4  * Uses 2 threads to populate frontier queues
5  */
6
7  //Necessary Libraries
8  #include <bits/stdc++.h>
9  #include <omp.h>
10 using namespace std;
11
12 //Defining Threads used
13 #define NUM_THREADS 4
14 #define POP_THREADS 2
15
16 //discovering for a specific queue and is shared among the threads
17 void discoverLevel(int &N,int type,vector<queue<pair<string,int>>>
    &q,vector<queue<pair<string,int>>> &tq,int &goal,vector<string>
    &p,vector<vector<int>> &G){
18     int l, r;
19     if(type % 2 == 0){
20         l = 1;
21         r = N;
22         r /= 2;
23     }
24     else{
25         l = N;
26         l /= 2;
27         l++;
28         r = N;
29     }
30     while(!q[type].empty()){
31         pair<string,int> node = q[type].front();

```

```

32     q[type].pop();
33     if(node.second == goal) p.push_back(node.first);
34     else{
35         int nodeidx = node.second;
36         string path = node.first;
37         for(int nextnode=1; nextnode<=r; nextnode++){
38             if(G[nodeidx][nextnode] == 1){
39                 string newpath = path + "->" + to_string(
nextnode);
40                 G[nodeidx][nextnode] = -1;
41                 G[nextnode][nodeidx] = -1;
42                 if(nextnode == goal){
43                     p.push_back(newpath);
44                     continue;
45                 }
46                 tq[type].push({newpath,nextnode});
47             }
48         }
49     }
50 }
51 return;
52 }
53
54 //random graph generator
55 void RandomGraph(int &N,vector<vector<int>> &G){
56     srand(time(NULL));
57     for(int i=1;i<=N;i++){
58         for(int j=i;j<=N;j++){
59             int edge = rand() % 2;
60             G[i][j] = edge;
61             G[j][i] = edge;
62         }
63     }
64     return;
65 }
66
67 //custom graph generator
68 void CustomGraph(vector<vector<int>> &G){
69     int u, v, edges;
70     cout << "Enter the number of edges: ";
71     cin >> edges;
72     for(int i=0;i<edges;i++){
73         cout << "Enter u & v for edge connection: ";
74         cin >> u;
75         cin >> v;
76         G[u][v] = 1;
77         G[v][u] = 1;
78     }
79     return;
80 }
81
82 int main(){
83     //Initializing parameters
84     int N, op, s, g;
85     cout << "Define the number of nodes in your Graph: ";
86     cin >> N;
87     cout << "Enter 1 to generate a random graph, 2 for your own

```

```

graph: ";
88 cin >> op;
89
90 //Graph Generation
91 vector<vector<int>> G(N+1,vector<int>(N+1,0));
92 if(op == 1) RandomGraph(N,G);
93 else CustomGraph(G);
94 vector<vector<int>> GCopy = G;
95
96 //Queues for the parallelism
97 vector<queue<pair<string,int>>> q(NUM_THREADS);
98 vector<queue<pair<string,int>>> tq(NUM_THREADS);
99 cout << "Enter root node: ";
100 cin >> s;
101 cout << "Enter goal node: ";
102 cin >> g;
103 if(s <= (N / 2)){
104     q[0].push({to_string(s),s});
105     q[1].push({to_string(s),s});
106 }
107 else{
108     q[2].push({to_string(s),s});
109     q[3].push({to_string(s),s});
110 }
111
112 //Running BFS in parallel
113 double start_time = omp_get_wtime();
114 vector<string> paths;
115 while((paths.size() == 0) && (!q[0].empty() || !q[1].empty() ||
    !q[2].empty() || !q[3].empty())){
116
117     //Running BFS Traversal in 4 parallel threads
118     omp_set_num_threads(NUM_THREADS);
119     #pragma omp parallel
120     {
121         int thread_id = omp_get_thread_num();
122         discoverLevel(N,thread_id,q,tq,g,paths,G);
123     }
124
125     //Populating frontier queues in 2 parallel threads
126     omp_set_num_threads(POP_THREADS);
127     #pragma omp parallel
128     {
129         int thread_id = omp_get_thread_num();
130         if(thread_id == 0){
131             while(!tq[0].empty()){
132                 q[0].push(tq[0].front());
133                 q[1].push(tq[0].front());
134                 tq[0].pop();
135             }
136             while(!tq[2].empty()){
137                 q[0].push(tq[2].front());
138                 q[1].push(tq[2].front());
139                 tq[2].pop();
140             }
141         }
142         else{

```

```

143         while(!tq[1].empty()){
144             q[2].push(tq[1].front());
145             q[3].push(tq[1].front());
146             tq[1].pop();
147         }
148         while(!tq[3].empty()){
149             q[2].push(tq[3].front());
150             q[3].push(tq[3].front());
151             tq[3].pop();
152         }
153     }
154 }
155
156 double end_time = omp_get_wtime();
157
158 //Printing Solutions and execution time of algorithm
159 cout << "----Parallelized BFS----\n";
160 for(auto sol : paths) cout << sol << "\n";
161 cout << "Computed in " << end_time - start_time << " units of
time\n";
162
163 //Non-Parallelized BFS
164 start_time = omp_get_wtime();
165 queue<pair<int,pair<string,int>>> Q;
166 Q.push({0,{to_string(s),s}});
167 vector<string> ans;
168 while(!Q.empty()){
169     pair<int,pair<string,int>> cur = Q.front();
170     Q.pop();
171     int nodeidx = cur.second.second, level = cur.first;
172     string path = cur.second.first;
173     if(nodeidx == g){
174         ans.push_back(path);
175         while(!Q.empty() && Q.front().first == level){
176             if(Q.front().second.second == g) ans.push_back(Q.
front().second.first);
177             Q.pop();
178         }
179         break;
180     }
181     else{
182         for(int i=1;i<=N;i++){
183             if(GCopy[nodeidx][i] == 1){
184                 GCopy[nodeidx][i] = -1;
185                 GCopy[i][nodeidx] = -1;
186                 string newpath = path + "->" + to_string(i);
187                 Q.push({level+1,{newpath,i}});
188             }
189         }
190     }
191 }
192 end_time = omp_get_wtime();
193
194 //Printing solutions and time taken for the non parallelized
BFS
195 cout << "\n----Normal BFS----\n";
196 for(auto p : ans) cout << p << "\n";

```


```

197     cout << "Computed in " << end_time - start_time << " units of
198         time\n";
199     return 0;
}


```

Results

- For less than or equal to 100 nodes sequential bfs outperforms parallelized bfs:




(a) Performance in seconds(upto 99 nodes)




(b) Performance in %(upto 99 nodes)

- When there are more than 100 nodes we can see parallelized bfs starts outperforming sequential bfs:




(c) Performance in seconds (100-110 nodes)




(d) Performance in %

- The code is then run for a graph with 200,500 and 1000 nodes and here are the results:




(e) Performance in seconds (200,500 and 1000 nodes)




(f) Performance in % (200,500 and 1000 nodes)

- Finally the code is run for a graph with 5000 nodes and an improvement of approximately 6500% is observed.



(g) Performance in seconds (upto 5000 nodes)



(h) Performance in % (upto 5000 nodes)